Synthesising Clustering Algorithms from Benchmark Failure Modes:

Basis Mixtures, Region-Adaptive Pipelines, Learned Routers and a Behaviour-Cloned Clustering Agent

Ashail Maharaj
https://github.com/Ashail33/clustbench

October 2026

Abstract

We report fifteen rounds of algorithm development carried out inside Clustbench [1], a benchmark harness that re-ranks every candidate against a registry of clustering algorithms on every run. The algorithms were not designed up front; each round started from a failure mode the previous round's benchmark exposed. We contribute (i) a family of exponential-family mixture models whose component log-densities are linear in a shared basis, with random Fourier (FMM), graph Laplacian (LMM) and autoencoder (AMM) bases; (ii) a physics-inspired method (MRIC) that clusters per-point relaxation signatures; (iii) five targeted one-line fixes of classical methods; (iv) three composed pipelines, RAPID, AURA and META-CLUSTERER, each iterated through a v1 $\rightarrow$ v2 $\rightarrow$ v3 cycle; (v) a graph-community method, `louvain_knn`, built to close a coverage gap; (vi) ten learned algorithm selectors (*routers*), including fingerprint-kNN, per-algorithm regressors, probe-based and stacked variants; (vii) four machine-proposed algorithms; and (viii) `rl_pipeline`, a behaviour-cloned policy that clusters by chaining eight primitive operators. All results come from a rerun of the benchmark with validity controls (five seeds, post-hoc tests, real-data provenance checks, timing of the algorithm call alone). LMM is the strongest single algorithm we built: rank 8 of 48 in-distribution (mean ARI 0.765) and rank 2 of 49 on an out-of-distribution split (0.858), statistically tied with spectral clustering but almost 50 times faster at n=30,000. The best routers beat it by 0.05 ARI in-distribution, but retrained without the dataset family under test they lose 0.03–0.07 and are statistically tied with it. Targeted fixes frequently traded the diagnosed failure for a new one. Hand-coded dispatch between versions helped only when the versions were complementary on detectable regimes. The RL agent learned to fire the spectral-embedding move from imitation, but not the multi-step sequences that complete it, and a published smoke result on circles did not reproduce with a retrained checkpoint.

1 Introduction

No clustering algorithm is best on all data [2, 3], and the classical families fail in characteristic ways. Centroid methods cannot represent non-convex clusters. Mean-based updates are dragged by outliers. Density methods depend on a scale parameter. Graph methods pay quadratic eigendecomposition costs. A companion paper [1] describes Clustbench, the harness used here. It runs a registry of algorithms on a grid of synthetic, real, time-series, graph and adversarial datasets, records quality and cost, and aggregates by average rank and the Friedman test [4].

This paper reports what we built with it. The process was empirical iteration: run the benchmark, name the narrowest failure mode visible in the per-shape and per-outlier breakdowns, and build the smallest change that addresses it. Section 2 makes this loop explicit. Section 3 describes the algorithms. Section 4 evaluates them in-distribution and on an out-of-distribution split, and Section 5 states the lessons, several of them negative.

2 Methodology: empirical iteration

The benchmark's per-algorithm analysis reduced almost every weakness to one of six failure modes: (F1) outlier sensitivity of mean-based updates; (F2) non-convex shape failure of Voronoi-cell methods; (F3) super-linear scaling; (F4) hyperparameter brittleness; (F5) local-minimum search; and (F6) a wrong structural assumption (e.g. subspace clustering on isotropic blobs). Each failure mode has a short menu of known remedies. That observation motivates a five-stage loop:

1. **v1, synthesis.** Compose the mechanisms that already won on each axis: k-means EM for convex blobs, a graph-Laplacian embedding for non-convex shape, posterior weighting for outliers, single-pass trees for scale, data-adaptive parameters for robustness, and diverse ensembles.
2. **Diagnose.** Benchmark v1 and name its *specific* new failure, e.g. "ARI 0.00 on circles because a full-covariance GMM absorbs both rings into one Gaussian."
3. **v2, targeted fix.** Change only the mechanism the diagnosis names.
4. **v3, dispatch.** If v1 and v2 win on different regimes, route between them on a cheap data signature.
5. **Learned dispatch.** Replace hand-coded routing with a model trained on the benchmark's own (task fingerprint, algorithm, ARI) table.

The loop stops when per-iteration lift falls within noise, when the router concentrates on one choice, or when no regime in the grid can show a lift. In that last case the right move is to enlarge the benchmark, not the algorithm.

3 Algorithms

3.1 Exponential-family basis mixtures: FMM, LMM, AMM

GMMs [5] assume Gaussian components. We replace each component with an exponential family whose sufficient statistics are a shared, fixed basis $\phi : \mathbb{R}^d \to \mathbb{R}^M$:

$$p(x) = \sum_{k=1}^{K} \pi_k \, p_k(x; \alpha_k), \qquad p_k(x; \alpha_k) \propto \exp(\alpha_k^\top \phi(x)).$$

The log-partition function is intractable for a general basis. We sidestep it by *self-normalising over the data*:

$$\log p_k(x_i) = \alpha_k^\top \phi(x_i) - \log \sum_{j=1}^{n} \exp(\alpha_k^\top \phi(x_j)),$$

which turns p_k into a discrete distribution over the observed points. The E-step computes responsibilities γ_{ik} as usual. The M-step is concave in each α_k and is solved with one damped Newton step,

$$(\mathrm{Cov}_{q_k}[\phi] + \lambda I)\, \Delta\alpha_k = \tfrac{1}{N_k} \textstyle\sum_i \gamma_{ik} \phi(x_i) - \mathbb{E}_{q_k}[\phi] - \lambda \alpha_k,$$

where q_k is the softmax of $\alpha_k^\top \phi(x_j)$ over the data. The Hessian is a covariance matrix, so EM converges in a handful of outer iterations, and every iteration is logged as a trajectory step. The three members differ only in ϕ:

- **FMM** uses random Fourier features [6], $\phi(x) = [\cos \omega_m^\top x, \sin \omega_m^\top x]_{m=1}^{M}$. The frequencies are drawn from a mixture of bandwidths centred on the inverse median pairwise distance, so one model resolves both fine and coarse structure.
- **LMM** uses the bottom eigenvectors ψ of the normalised kNN-graph Laplacian [7], i.e. the spectral embedding, weighted by a per-cluster heat kernel $W_m(\tau_k) = e^{-\tau_k \lambda_m / 2}$. Here λ_m is the m-th Laplacian eigenvalue, so τ_k is a learned graph-diffusion bandwidth: small τ_k keeps neighbourhood detail, large τ_k averages over a connected component. LMM can be read as spectral clustering with soft, EM-trained assignments in place of a final k-means.

- **AMM** uses the bottleneck activations of a shallow autoencoder. It is intended for high-dimensional sparse inputs, where kNN distances lose meaning.

3.2 MRIC: magnetic-resonance-inspired clustering

MRI separates tissues because nuclei in different chemical environments relax at different rates. MRIC treats each point as a spin and its kNN neighbourhood as its environment. The first principal axis plays the static field B_0. Local density (inverse mean kNN distance) and local anisotropy (from the eigenvalues of the local covariance) map to per-point T_1/T_2 relaxation times. A simulated free-induction-decay signal is sampled into a per-point signature, and k-means clusters the signatures.

3.3 Targeted fixes of classical methods

Each fix changes one mechanism named by the diagnosis:
- `dbscan_auto`: ε set from the knee of the k-distance plot [8] (F4).
- `meanshift_robust`: bandwidth estimated on a trimmed sample (F1/F4).
- `clarans_pp`: k-means++ medoid seeding [9] (F5).
- `pwcc_diverse`: PWCC [10] with heterogeneous bases $\{k\text{-means}, \text{spectral}, \text{GMM}\}$ in place of three k-means variants (F2).
- `kmeans_trimmed`: drop the farthest 10% of members from each centroid update (F1).

3.4 Composed pipelines and their v1 → v3 lineages

RAPID (Region-Adaptive Partitioning with Iterative Density). *v1* partitions the data into density-connected regions with auto-ε DBSCAN. It computes a convexity signature per region and sends convex regions to k-means++ and non-convex regions to spectral clustering, then reassigns DBSCAN noise by 1-NN. Diagnosis: the k-distance knee is contaminated by the outliers it should ignore (outlier ΔARI -0.145). *v2* adds a stage-0 LOF [11] pre-filter that holds out the top 10% of LOF scores. *v3* runs the v2 path only when the estimated LOF outlier fraction exceeds 5%.

AURA (Adaptive Unified Robust Algorithm). *v1* computes a Nyström-approximated [12, 13] normalised-Laplacian embedding Y and fits a full-covariance GMM on it, aiming for spectral shape-handling plus GMM outlier robustness. Diagnosis: ARI 0.00 on circles, because the GMM merges both rings into one elongated component and the eigen-columns differ in scale by $\sim$20×. *v2* z-scores Y column-wise and runs k-means. *v3* counts the informative embedding axes, $r = |\{j : \mathrm{sd}(Y_{:j}) \geq 0.1 \max_{j'} \mathrm{sd}(Y_{:j'})\}|$, and routes $r \geq 2$ to the v2 path (z-scored k-means) and $r < 2$ to the v1 path (raw GMM). Moons have one informative axis and circles two, so each version runs on its own regime.

META-CLUSTERER. *v1* computes four fingerprints (PCA effective dimension, Laplacian eigengap, LOF outlier fraction, kNN density skew) and dispatches with hand-written rules. Diagnosis: the eigengap rule never fires on clean circles, whose gap is *large*. *v2* replaces the eigengap with a convexity ratio of the k-means partition and the kNN-graph modularity [14] of that partition. It also adds a silhouette probe that fits three candidates on a 20% subsample when a rule fires with low margin. *v3* runs both routers and probes only when they disagree.

3.5 `louvain_knn`: a graph-native method for a coverage gap

Enlarging the benchmark to graph-node features exposed Zachary's karate club [15] as a regime where every algorithm scored below 0.43 ARI. With 34 nodes and 16 features, re-embedding the features drowns the Fiedler signal. `louvain_knn` builds a symmetrised kNN graph on the

Table 1: Learned router generations.

version	mechanism
v1	7 features ($\log n$, d, k, PCA effective dimension, convexity CV, LOF outlier fraction, density skew); K=5 nearest training tasks; pick the lowest mean ARI rank.
v2	15 features: v1 plus 3-iteration k-means++ probe statistics (inertia ratio, silhouette, size CV, centroid-shift ratio), Levina–Bickel intrinsic dimension [17], multi-scale kNN distances.
v3	17 features: v2 plus an auto-ε DBSCAN probe (clusters found, noise fraction); one distance-weighted kNN *regressor per algorithm* predicting ARI; dispatch on the argmax.
v4	20 features (adds real-data, image-domain and outlier-extremity flags); gradient-boosted regressors per algorithm; fall back to v3 if the two disagree by > 0.05 ARI.
v5	Stacker choosing between v3 and v4 per task [18].
v6	Probe router: run candidates on a subsample and dispatch on silhouette.
v6b	v6 with a multi-metric probe (silhouette, Calinski–Harabasz, Davies–Bouldin, k-neighbour purity).
v6c	v6b with an inverse-cost term at 20% weight.
v7	Random-forest stacker over {v3, v6}, trained to predict where v6 beats v3 by > 0.02.
v8	11-feature enriched fingerprint (distance kurtosis, mass Gini, kNN density variance, covariance effective rank), out-of-distribution-triggered top-3 ensemble, aggregation by mean ARI.

features and optimises modularity directly with Louvain [16]. It then reconciles the discovered community count with the requested k: it merges the smallest community into its most densely connected neighbour, or splits the largest community by spectral bipartition. It falls back to a self-contained NumPy Louvain implementation, and then to spectral clustering, if dependencies are missing.

3.6 Learned routers

A router computes a fingerprint of the input, chooses an algorithm from the registry, and runs it. It is trained on the benchmark's own results: every historical task is regenerated from its specification, fingerprinted, and joined with every algorithm's ARI on it. Training happens at import time, so a router retrains whenever the results file changes. Table 1 summarises the ten generations.

All routers exclude themselves and their siblings from the candidate set to prevent recursion. For honest in-distribution evaluation they drop an exact fingerprint self-match (task-level leave-one-out).

3.7 Machine-proposed algorithms

Card mutator. Clustbench encodes every algorithm as a card listing its mechanisms, inductive biases and complexity [1]. A mutator searches card space with swap, add-bias, crossover and stack operators and scores candidates by a bias-matching ARI upper bound. Its top proposal, `mutant_kmeans_meta`, crosses k-means with META-CLUSTERER-v2. It keeps the k-means partition when a convexity probe and a silhouette check both pass, and otherwise hands off to the META-v2 router.

Multi-agent discovery (round 14). LLM agents running in parallel proposed algorithms, ordinary datasets and adversarial datasets. Separate agents implemented, integrated and

benchmarked them. Three algorithms came out of this round:

- `density_peaks_bridge`: density peaks [19] plus one round of normalised-Laplacian label spreading [20], which reassigns borderline halo points and keeps true outliers.
- `bico_stream`: a BICO-style [21] streaming coreset with merge-and-reduce [22], followed by weighted k-means++ on the coreset.
- `isomap_bgmm`: an Isomap [23] geodesic embedding followed by a Dirichlet-process Bayesian GMM [24].

Trident (round 15). Designed from the round-14 winners table, Trident runs one cheap probe per winning family (centroid: k-means; mixture: diagonal GMM; graph: `louvain_knn`). It gates on silhouette and pairwise-ARI agreement between the probes, and escalates to `isomap_bgmm` on a heavy-tail signature or to AURA-v3 on density variation. Selection is built into the algorithm rather than delegated to an external router.

3.8 `rl_pipeline`: clustering as sequential decision-making

Clustbench's trajectory layer records iterative algorithms as state–action–cost sequences [1]. We ask whether a policy can *compose* an algorithm from primitives. The environment has eight actions drawn from the registry's dominant mechanisms: `kpp_init`, `assign_to_centers`, `update_centers` (k-means); `medoid_swap` (CLARANS); `ward_merge` (Ward); `density_partition` (DBSCAN); `eigen_embed` (spectral); and `outlier_trim`. The state is a 20-dimensional vector: n, d, k, cost, silhouette, plateau signals, outlier mask statistics and centroid statistics. The reward is $\mathrm{clip}(-\Delta c_t / c_0)$.

Hand-written action sequences reproduce their source algorithms: a k-means sequence reaches ARI 1.0 on blobs, and the spectral sequence `kpp`→`eigen_embed`→`assign`→`update` reaches 1.0 on circles. For behaviour cloning [25], we drove k-means, CLARANS, Ward, spectral clustering and DBSCAN through the environment, which produced 50 episodes and 419 (state, action, reward) rows. We trained an MLP policy (cross-entropy, with count$^{-1/4}$ class re-weighting) and an MLP value head (discounted return, γ=0.95) on an 80/20 episode split. At inference the agent runs 16 stochastic rollouts, and a tiered selector keeps the best non-embedding rollout with silhouette ≥ 0.40, or else the best embedding rollout.

4 Results

Unless stated otherwise, numbers come from the controlled benchmark rerun [1]:

- an in-distribution split of 48 algorithms on 140 tasks (28 task variants from 20 dataset families, seeds 1–5, $n \leq 1000$);
- a leave-one-dataset-family-out (LOFO) evaluation of the routers on the same tasks;
- an out-of-distribution (OOD) split of 49 algorithms on 35 tasks with unseen (n, d, k), outlier extremities and seeds 100–104;
- a large-n sweep up to $n = 50{,}000$.

Only `fit_predict` is timed (median of three runs, one BLAS thread). Lineage tables (Tables 2–3) come from the round in which each version was introduced, on the 16-task demonstration grid and before the controls, and are marked as such.

4.1 Targeted fixes

Table 2 reports the five fixes. Auto-ε is the single largest intervention in the project (+0.589): it takes DBSCAN from one cluster on every task to ARI 0.83 on circles. Three fixes exposed second-order failures.

- `kmeans_trimmed` barely moves, because a 10% trim cannot absorb a 20% outlier rate.

Table 2: Targeted fixes (16-task demonstration grid, round 2).

fix	ARI before	ARI after	ΔARI	wall time (s)
`dbscan` → `dbscan_auto`	0.000	0.589	+0.589	0.020 → 0.035
`meanshift` → `meanshift_robust`	0.470	0.622	+0.151	0.229 → 0.288
`clarans` → `clarans_pp`	0.581	0.656	+0.075	0.018 → 0.017
`pwcc` → `pwcc_diverse`	0.702	0.751	+0.049	0.145 → 0.223
`kmeans` → `kmeans_trimmed`	0.701	0.703	+0.003	0.024 → 0.024

Table 3: v1/v2/v3 lineages (16-task demonstration grid, rounds 3–4). Per-shape columns are mean ARI. Rank is out of 31.

pipeline	version	mean ARI	mdcgen	aniso.	moons	circles	rank
AURA	v1	0.733	0.87	1.00	0.37	0.00	9
	v2	0.764	0.77	1.00	0.01	1.00	–
	v3	**0.875**	0.86	1.00	0.55	1.00	**2**
META	v1	0.774	0.90	1.00	0.42	0.16	–
	v2	**0.879**	0.90	1.00	0.42	1.00	4
	v3	0.827	0.90	1.00	0.42	0.58	8
RAPID	v1	**0.843**	0.76	1.00	0.85	0.84	5
	v2	0.774	0.73	0.91	0.85	0.57	–
	v3	0.843	0.77	1.00	0.85	0.84	7

- `clarans_pp` is *more* outlier-sensitive than CLARANS, because k-means++ samples proportionally to squared distance and so prefers outliers as seeds.
- `dbscan_auto`'s knee is itself inflated by outliers (outlier ΔARI −0.475).

4.2 Pipeline lineages

Table 3 follows each pipeline through v1, v2 and v3.

- **AURA** is the clean success of meta-dispatch. v2 traded moons (0.37 → 0.01) for circles (0.00 → 1.00), because z-scoring amplifies noise in the near-empty eigen-directions of moons. v3's effective-rank switch recovered both (moons 0.55, circles 1.00) and climbed from rank 9 to rank 2 of 31, with +0.142 ARI over v1.
- **META-CLUSTERER** v2 was a strict improvement: replacing a brittle rule fixed circles (0.16 → 1.00) and broke nothing. v3's probe then *lowered* the score by overruling correct v2 routes, because subsample silhouette does not track ARI on non-convex data.
- **RAPID** v2 closed most of the outlier gap (−0.145 → −0.050) but regressed on shape, because LOF removes real points when cluster density is comparable to outlier density. v2 never beat v1 on any regime, so v3's gate always selected v1 and v3 tied it.

RAPID v1 was, at the time, the only algorithm in the registry strong on both moons (0.85) and circles (0.84).

4.3 Final standings

Table 4 places every new algorithm in the controlled rerun, and Figure 1 breaks the results down by dataset family. Configuration-level 95% bootstrap intervals are about ±0.09 ARI, and a Nemenyi critical difference cannot separate the top eleven algorithms in-distribution. Read adjacent ranks as ties unless a pairwise test says otherwise [1]. Four results stand out.

- **LMM is the best single algorithm we built.** It ranks 8th of 48 in-distribution, behind only the routers, and 2nd of 49 OOD (0.858, against 0.860 for spectral clustering; the difference is not significant). It solves circles (1.00), and on the 15 contaminated tasks its inlier-only

Table 4: Controlled rerun. Mean ARI with rank by mean ARI in-distribution (48 algorithms) and OOD (49 algorithms); LOFO is the routers' mean ARI when retrained without the family under test. Median wall time (`fit_predict` only) is in-distribution. Classical baselines for reference.

algorithm	ARI (in)	rank	LOFO	ARI (OOD)	rank	wall (s)
`lmm`	0.765	8	–	0.858	2	0.033
`fmm`	0.657	28	–	0.728	28	0.152
`mri`	0.562	40	–	0.617	41	0.032
`louvain_knn`	0.714	13	–	0.846	3	0.098
`aura_v3`	0.737	10	–	0.831	8	0.024
`aura_v2`	0.714	14	–	0.832	7	0.040
`rapid`	0.702	16	–	0.727	29	0.051
`rapid_v3`	0.699	21	–	0.726	30	0.065
`meta_clusterer_v2`	0.701	18	–	0.628	39	0.084
`pwcc_diverse`	0.701	17	–	0.756	19	0.069
`mutant_kmeans_meta`	0.719	12	–	0.756	18	0.058
`trident`	0.725	11	–	0.826	10	0.148
`bico_stream`	0.677	25	–	0.719	32	0.030
`isomap_bgmm`	0.634	34	–	0.766	16	0.160
`density_peaks_bridge`	0.546	41	–	0.708	34	0.035
`rl_pipeline`	0.604	37	–	0.764	17	2.434
`learned_router` (v1)	0.788	7	0.728	0.808	11	0.203
`learned_router_v2`	0.792	6	0.747	0.789	13	0.144
`learned_router_v3`	**0.815**	**1**	0.745	0.845	4	0.109
`learned_router_v4`	0.814	4	**0.784**	0.831	9	0.235
`learned_router_v7`	0.815	3	0.745	0.845	6	0.167
`learned_router_v8`	0.808	5	0.748	0.800	12	0.103
`spectral`	0.763	9	–	**0.860**	**1**	0.027
`gmm`	0.699	20	–	0.774	15	0.014
`agglomerative` (Ward)	0.685	23	–	0.754	20	0.006
`kmeans`	0.655	30	–	0.731	23	0.006

ARI is 0.990. It also scales far better than spectral clustering on the same graph Laplacian: a log–log slope of 1.5 against 2.9, and 6 s against 297 s at n=30,000. Of all the mechanisms we measured, the graph-Laplacian basis is the most universal.

- **`louvain_knn` closed the gap it was built for and generalises.** It is the best non-router on the karate graph (0.73), moons (1.00) and Digits (0.55), ranks 3rd OOD (0.846), and keeps moons at 1.00 up to n=50,000. It is weak on imbalanced blobs (0.57) and power-law Student-t mixtures (0.50), where Louvain merges or splits communities by size.
- **AURA-v2/v3 generalise better than they rank.** They rank 14th and 10th in-distribution but 7th and 8th OOD (0.832 and 0.831), and AURA-v3's Nyström path is near-linear (slope 1.1).
- **FMM, MRIC and the round-14 proposals are mid-pack or lower.** Density peaks with a bridge won its own mini-grid but ranks 41st of 48 overall.

4.4 Routers: in-distribution gains that do not transfer

Figure 2 plots every router generation on three evaluations.

- **In-distribution, routers hold the top seven ranks.** The best beats LMM by 0.051 ARI (paired Wilcoxon, Holm $p = 2 \times 10^{-4}$).
- **Within the lineage, v3 beats v1** (Holm $p = 0.003$). v2's probe features and v3's switch from rank voting to per-algorithm ARI regression each added roughly 0.02–0.03 ARI when introduced.

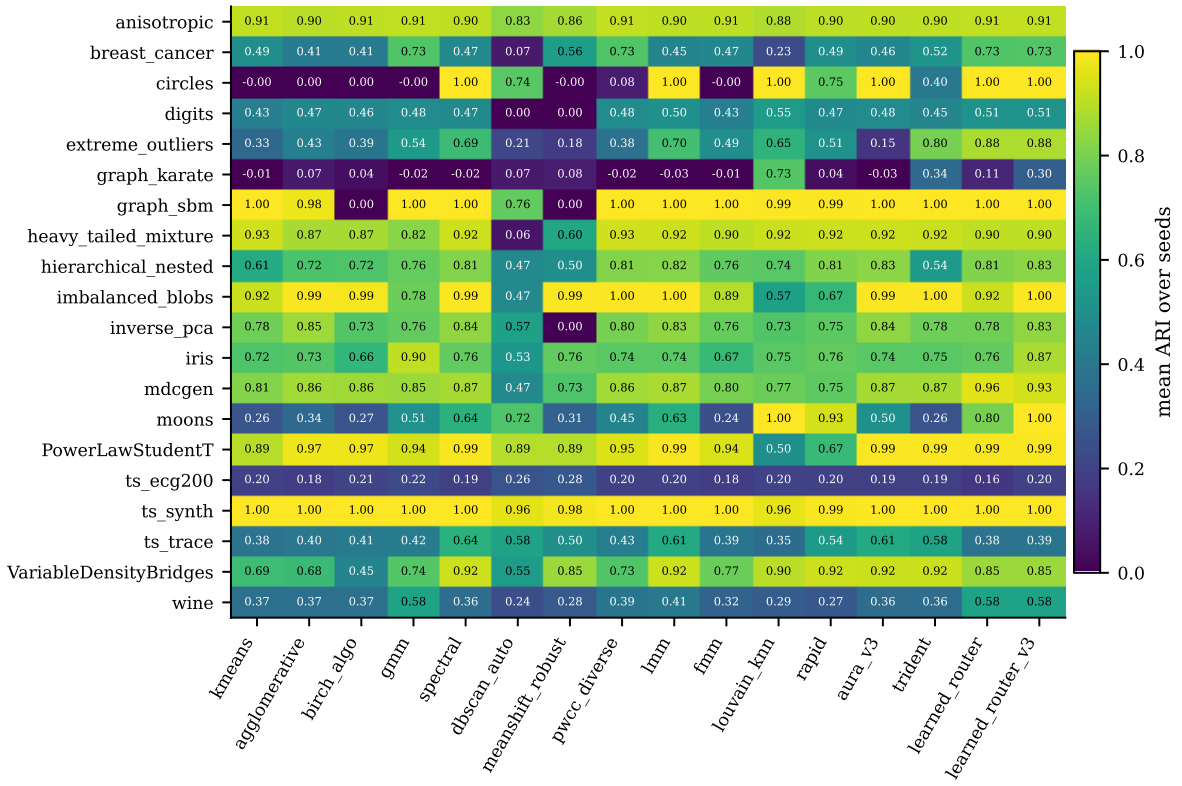


Figure 1: Mean ARI by dataset family (in-distribution) for classical baselines and the new algorithms. `louvain_knn` is the only non-router to solve moons and the karate graph together; LMM, spectral clustering and AURA-v3 solve circles; GMM is strongest on Iris and Wine; the real UCR time series (ECG200, Trace) are hard for every algorithm.

- **Architectural changes helped only when they added a new signal.** v4 used the same features with a different model and overrode v3 on 51 of 52 tasks without improving, so the v5 stacker over {v3, v4} degenerated to always-v3. v6's probe signal is different in kind and won exactly one task v3 missed (inverse-PCA, 0.00 → 0.38). That was enough for the v7 stacker over {v3, v6} to take rank 1 in round 8.
- **Silhouette is a biased probe.** As a probe metric it favours compact wrong answers on non-convex data, and adding a 20% cost term (v6c) dethroned spectral clustering on close calls where it was correct.

Leave one family out. Retraining each router without the family under test removes the leakage of sibling seeds.

- v3, v5 and v7 fall from 0.815 to 0.745, v1 from 0.788 to 0.728, and v8 from 0.808 to 0.748.
- v4, whose gradient-boosted regressors interpolate more smoothly, is the most robust: 0.814 to 0.784.
- The losses concentrate where a family's winner is unusual. With moons held out, v3 never learns to send moons-like data to `louvain_knn` and loses 0.70 ARI there; the karate graph costs 0.32.
- Paired against LMM and spectral clustering on the same 140 tasks, no LOFO router differs significantly (Holm $p \geq 0.07$; mean differences from -0.037 to $+0.020$).

Unseen configurations. On the OOD split, v3, v5 and v7 (0.845) are level with spectral clustering (0.860), LMM (0.858) and `louvain_knn` (0.846); none of these differences survives

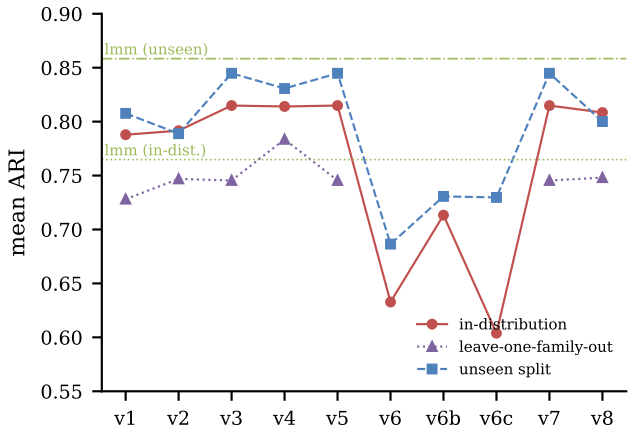


Figure 2: Mean ARI of each router generation in-distribution (task-level leave-one-out), leave-one-family-out (stacked and probe routers v5–v6c not evaluated) and on the OOD split. LMM is shown for reference.

Holm correction. Typical remaining failures are inverse-PCA at d=150 (v3 0.72 against 0.85 for Ward) and outlier extremity 10 (v3 0.74 against 0.92 for `rl_pipeline`). v8's richer fingerprint and OOD-triggered ensemble again did not help (0.800).

An earlier version of this paper reported that v3–v7 lost 0.11–0.13 ARI on the OOD split and that v1 generalised best. That result came from routers trained on a narrower grid (three seeds, no inverse-PCA, synthetic stand-ins for the UCR series) and does not reproduce. The robust conclusion is the LOFO one: a router's advantage comes from having seen a family's winner.

Self-extension (development history, before the controls). Routers retrain from the results file, so adding `louvain_knn` and re-running the benchmark was enough for every fingerprint router to start dispatching to it on the karate graph. This raised the mean ARI of v3, v4, v5 and v7 by 0.029–0.038, the largest router gain since v3, with zero router code changed. Round 15's additions likewise lifted v7 from 0.871 to 0.912. The LOFO result above is the other side of the same mechanism: a router can only dispatch to a winner it has seen.

4.5 Machine-proposed algorithms

The card mutator predicted 0.926 ARI for `mutant_kmeans_meta`; it measured 0.747 at introduction and 0.719 in the controlled rerun (rank 12). It beats its parent `kmeans_trimmed` (0.654) by 0.065, so the hybrid is a real improvement, but the predicted magnitude was not. Bias matching can order architectures but cannot calibrate them.

Of the round-14 proposals, the two adversarial datasets did what they were designed to do: routers trailed non-routers by 0.036 and 0.068 ARI on them. Of the three new algorithms, none reached the top 20.

Trident (rank 11, 0.725) shows the limit of silhouette gating. On moons it scored 0.262 even though its own `louvain_knn` probe scored 0.997, because a compact wrong split can have a higher silhouette than a correct non-convex one.

4.6 `rl_pipeline`

The behaviour-cloned policy reached 0.59 validation accuracy at predicting the next action. Chance is 0.125 and the majority-class baseline is 0.41. In the original 300-point smoke test it scored ARI 1.00 on blobs and 1.00 on circles, where k-means scores −0.003. On circles it chose `eigen_embed` early: from imitation alone, it learned to switch to the spectral embedding when

the cost stops falling. On moons it scored 0.26, because a k-means-like rollout won the silhouette tie-break.

Reproducibility. The traces and checkpoint lived in a git-ignored directory, which is why the agent produced no rows in round 14. The agent now rebuilds both deterministically when they are missing. The rebuilt traces match the originals exactly: 419 rows, 50 episodes, and the same action counts. A checkpoint retrained on them (PyTorch 2.14) reproduces blobs (1.00) and moons (0.30), but scores *0.16* on the circles smoke test, not 1.00. That headline result therefore depended on one lucky training run.

Controlled rerun. The agent ranks 37th of 48 in-distribution (mean ARI 0.604) and is the slowest method (median 2.4 s per task), with a bimodal profile: 0.989 on the SBM graph and 0.907 on anisotropic blobs, but −0.001 on circles and 0.27 on moons at n=500. It does better OOD (rank 17, 0.764), where it is the best algorithm on outlier extremity 10 (0.92).

- At n=500 the policy still fires `eigen_embed`, but every rollout then stalls on `update_centers` no-ops. 419 imitation rows cannot teach the coordinated sequence that completes a spectral partition.
- This is the behaviour-cloning ceiling. Imitation of per-step decisions does not value a move by where the episode ends, and compounding errors off the demonstrator distribution are expected [26].

5 Lessons

1. **Targeted fixes often open new failure modes.** Two of the five one-line fixes (`dbscan_auto`, `clarans_pp`) and two of the three pipeline v2s (AURA, RAPID) traded the diagnosed failure for a new one. The clean pipeline win was META v2, where a brittle rule was strictly replaced.
2. **Dispatch pays only when versions are complementary on detectable regimes.** AURA-v3 (+0.142) succeeded because effective rank separates moons from circles. RAPID-v3 could not help because v2 never won. META-v3 lost because its probe was noisier than the v2 it refined.
3. **Stacking needs architectural diversity.** Re-fitting the same features with a different model (v4) added nothing. A different signal (v6's probe) added a real, if small, gain.
4. **Learned selection does not transfer to unseen families.** Routers beat every single algorithm in-distribution, but with the test family held out they are statistically tied with LMM and spectral clustering. Only a family-level split shows this; a seed-level or configuration-level hold-out does not. Outside the benchmark's families, LMM, spectral clustering, `louvain_knn` and AURA-v2/v3 are as good a default and much cheaper.
5. **Grow the registry, not the router.** The biggest router gains came from adding algorithms and data, not from router redesigns.
6. **Internal-validity gates are shape-biased.** Silhouette-based probes (v6, META-v3, Trident) repeatedly pick compact wrong answers on non-convex data. A connectivity-aware criterion or shape prior is needed.
7. **Imitation recovers moves, not plans.** The RL agent learned to fire `eigen_embed` as a conditional move but not the sequence that makes it pay off, and its best smoke result depended on one training run.
8. **Silent failures bias means.** GMM's singular covariances, a YAML float parsed as a string, a missing default argument and git-ignored training files each removed rows from the benchmark without stopping it. Each missing row changed the means it was dropped from until it was root-caused.

Limitations. Most tasks are small ($n \leq 1000$; the scaling sweep reaches $n = 50{,}000$ on two families only). Twenty families give wide intervals, so many adjacent ranks are ties. Lineage tables come from earlier, smaller grids run before the validity controls. Algorithms run with fixed hyperparameters. Many ideas here have prior art: algorithm selection [27], meta-learning for clustering [28, 29], landmarking [30] (our probe features) and stacked generalisation [18]. Our contribution is their systematic, re-runnable evaluation against a broad registry, the basis-mixture family, and the trajectory-trained agent.

6 Conclusion and future work

Starting from benchmark failure modes, we built 25 algorithms and ten learned selectors and evaluated all of them against one another, including on data none had seen. The most durable results are a single algorithm (LMM, an EM-trained mixture over a graph-Laplacian basis, accurate and near-linear in practice) and a negative result: learned routers win only on dataset families they have seen. Three directions follow:

- fine-tune the behaviour-cloned agent with PPO [31], rewarding terminal partition quality, using the trajectory corpus every benchmark run now produces;
- replace specific weak components, such as CLARANS's swap proposal and Chameleon's merge rule, with models trained on trajectory data;
- give selection gates a shape-aware criterion, e.g. penalise silhouette when the number of kNN-graph components disagrees with k.

Availability. All algorithms, configurations, results and the figure scripts are at `https://github.com/Ashail33/clustbench`.

References

[1] Ashail Maharaj. Clustbench: A reproducible, trajectory-instrumented, self-extending benchmark for clustering algorithms, 2026. Companion paper, `https://github.com/Ashail33/clustbench/tree/master/paper/clustbench`.

[2] Jon Kleinberg. An impossibility theorem for clustering. In *Advances in Neural Information Processing Systems*, volume 15, 2002.

[3] David H. Wolpert and William G. Macready. No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1):67–82, 1997.

[4] Janez Demšar. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7:1–30, 2006.

[5] A. P. Dempster, N. M. Laird, and D. B. Rubin. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society: Series B*, 39(1):1–38, 1977.

[6] Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. In *Advances in Neural Information Processing Systems*, volume 20, 2007.

[7] Ulrike von Luxburg. A tutorial on spectral clustering. *Statistics and Computing*, 17(4): 395–416, 2007.

[8] Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. A density-based algorithm for discovering clusters in large spatial databases with noise. In *Proceedings of the 2nd International Conference on Knowledge Discovery and Data Mining*, pages 226–231, 1996.

[9] David Arthur and Sergei Vassilvitskii. k-means++: The advantages of careful seeding. In *Proceedings of the 18th Annual ACM-SIAM Symposium on Discrete Algorithms*, pages 1027–1035, 2007.

[10] Rasim M. Alguliyev, Ramiz M. Aliguliyev, and Lyudmila V. Sukhostat. Weighted consensus clustering and its application to big data. *Expert Systems with Applications*, 150:113294, 2020.

[11] Markus M. Breunig, Hans-Peter Kriegel, Raymond T. Ng, and Jörg Sander. LOF: Identifying density-based local outliers. In *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data*, pages 93–104, 2000.

[12] Christopher K. I. Williams and Matthias Seeger. Using the Nyström method to speed up kernel machines. In *Advances in Neural Information Processing Systems*, volume 13, 2001.

[13] Charless Fowlkes, Serge Belongie, Fan Chung, and Jitendra Malik. Spectral grouping using the Nyström method. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 26 (2):214–225, 2004.

[14] M. E. J. Newman. Modularity and community structure in networks. *Proceedings of the National Academy of Sciences*, 103(23):8577–8582, 2006.

[15] Wayne W. Zachary. An information flow model for conflict and fission in small groups. *Journal of Anthropological Research*, 33(4):452–473, 1977.

[16] Vincent D. Blondel, Jean-Loup Guillaume, Renaud Lambiotte, and Etienne Lefebvre. Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment*, 2008(10):P10008, 2008.

[17] Elizaveta Levina and Peter J. Bickel. Maximum likelihood estimation of intrinsic dimension. In *Advances in Neural Information Processing Systems*, volume 17, 2005.

[18] David H. Wolpert. Stacked generalization. *Neural Networks*, 5(2):241–259, 1992.

[19] Alex Rodriguez and Alessandro Laio. Clustering by fast search and find of density peaks. *Science*, 344(6191):1492–1496, 2014.

[20] Dengyong Zhou, Olivier Bousquet, Thomas Navin Lal, Jason Weston, and Bernhard Schölkopf. Learning with local and global consistency. In *Advances in Neural Information Processing Systems*, volume 16, 2004.

[21] Hendrik Fichtenberger, Marc Gillé, Melanie Schmidt, Chris Schwiegelshohn, and Christian Sohler. BICO: BIRCH meets coresets for k-means clustering. In *Algorithms – ESA 2013*, volume 8125 of *Lecture Notes in Computer Science*, pages 481–492, 2013.

[22] Sariel Har-Peled and Soham Mazumdar. On coresets for k-means and k-median clustering. In *Proceedings of the 36th Annual ACM Symposium on Theory of Computing*, pages 291–300, 2004.

[23] Joshua B. Tenenbaum, Vin de Silva, and John C. Langford. A global geometric framework for nonlinear dimensionality reduction. *Science*, 290(5500):2319–2323, 2000.

[24] David M. Blei and Michael I. Jordan. Variational inference for Dirichlet process mixtures. *Bayesian Analysis*, 1(1):121–143, 2006.

[25] Dean A. Pomerleau. ALVINN: An autonomous land vehicle in a neural network. In *Advances in Neural Information Processing Systems*, volume 1, 1989.

[26] Stéphane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics*, pages 627–635, 2011.

[27] John R. Rice. The algorithm selection problem. *Advances in Computers*, 15:65–118, 1976.

[28] Marcilio C. P. de Souto, Ricardo B. C. Prudêncio, Rodrigo G. F. Soares, Daniel S. A. de Araujo, Ivan G. Costa, Teresa B. Ludermir, and Alexander Schliep. Ranking and selecting clustering algorithms using a meta-learning approach. In *2008 IEEE International Joint Conference on Neural Networks*, pages 3729–3735, 2008.

[29] Daniel G. Ferrari and Leandro Nunes de Castro. Clustering algorithm selection by meta-learning systems: A new distance-based problem characterization and ranking combination methods. *Information Sciences*, 301:181–194, 2015.

[30] Bernhard Pfahringer, Hilan Bensusan, and Christophe Giraud-Carrier. Meta-learning by landmarking various learning algorithms. In *Proceedings of the 17th International Conference on Machine Learning*, pages 743–750, 2000.

[31] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.